Issues in Pragmatics

PLIN3001 – PLING204

Jonny McIntosh
jonathon.mcintosh@ucl.ac.uk

Lecture 9: Compositionality and Indexicalism

Overview

Last week, we looked at issues surrounding **compositionality**, the principle that the meaning (encoded meaning, truth conditional content, or etc.) of a complex expression is determined by the meanings of its parts and the way those parts are put together. We saw that there is some reason to think that language is *not* compositional, at least at the level of truth conditional content. For many of the examples we've looked at in previous weeks seem to show that the content of complex expressions are generally *not* determined by the contents of their parts and the way those parts are put together.

This week, I want to look at a challenge Jason Stanley has raised for those who would accept this result. Put briefly, Stanley thinks that theories on which language is not compositional in the relevant sense *over-generate*, i.e. they predict that certain interpretations to arise that in fact cannot arise. Stanley's own view is that language *is* compositional, and that apparent counter-examples can be accounted for by covert linguistic structure. On this view, called **indexicalism**, what is said varies much as the contextualist suggests, but the variation is the result of linguistically controlled processing triggered by **hidden indexical** that occurs in the sentence's logical form.

1 Contextualism and Compositionality

Before looking at Stanley, though, it is worth briefly summarising the case against the compositionality of language—and considering a response that Fiona suggested in questions. As I mentioned, Fiona's response is essentially the one that Recanati offers, and it's useful to consider it in a little detail before we think about Stanley's challenge.

The strongest case against compositionality revolves around cases like the following:

(1) 'It's raining' [*location constituent*]

(2) 'This leaf is green' [*painted/naturally green*]

The case can be put schematically as follows: holding constant the contents of the parts of these sentences, we can find pairs of contexts in which the sentences themselves have different contents (truth conditions). So the contents of these sentences cannot be determined by the contents of their parts and the ways that they are put together.

For example, the contents of the various parts of (2) seem to provide three things: a contextually determined leaf, which is the content of 'this leaf'; a contextually determined time, which is the content of the present tense of 'is'; and the property of being green, which is the content of 'green'. Nevertheless, we can find pairs of contexts in which the contents of these parts are all the same and the sentence's truth conditions vary.

2 Recanati's Solution

In effect, what Fiona pointed out at the end of last week's lecture was that the argument against compositionality only goes because we have assumed that the content of an expression is linguistically controlled—that the content of 'green' is the property of being green. Could we perhaps hang on to compositionality if we abandoned and took the content of an expression instead to be its *modulated* meaning? In the case of (2), we could then say that what enters composition is not the property of being green, but rather something different in the different contexts—perhaps the property of being painted green in one context and the property of being naturally green in others.

This is essentially Recanati's proposal. And it's not immediately clear what, if anything, could be wrong with it. To some extent, it parallels the standard understanding of indexicals: We don't compose the *encoded* meanings of the parts of, say, 'I am hungry' to get the result, THE SPEAKER IS HUNGRY. Rather, what goes into composition is the context-specific content of 'I', e.g. JONNY$_x$, which gives as a result, JONNY$_x$ IS HUNGRY. Similarly, what goes into composition, in Recanati's view, is not the *unmodulated* meanings of the parts of sentences, but rather the context-specific modulated meanings, like GREEN$_{\text{painted}}$. Of course, there's also a difference: context-specific meanings of indexicals are linguistically controlled, while context-specific modulated meanings are not. But it is not obvious this difference should be thought to be of significance.

Recanati seems to think it isn't significant. He argues that, strictly speaking, he only secures a relatively weak form of compositionality. For while the modulated meaning of a complex expression *depends on* the modulated meanings of its parts and the way they are put together, it won't generally be *determined* by them. For it may be the result of taking the result of composing the modulated meanings of its parts and the way they are put together and subjecting *that* to further modulation.[1] Recanati doesn't think

[1]This might be the story Recanati would tell about a case in which (1) is used to say that it is raining *heavily* in *London*. By modulating the meaning of 'raining', we get the concept RAIN* (i.e. an ad hoc

this is a problem, citing Peter Pagin (2005, p. 313): "having the context itself as an extra argument cannot be objected to as violating compositionality, since the meaning function takes a context argument to begin with". The thought here seems to be: it's OK to take context into account somewhere else, so it's OK to take into account here.

3 Some Worries

But is it OK? Here are two worries.

a. One worry relates to semantics, taken to be the study of the tacit knowledge employed in pairing expressions with their meanings—and in particular, their encoded meanings. Last week, I suggested that semantics, so construed, had better be compositional, since the capacity to pair expressions with their meanings is productive and systematic. Now, there is no problem *per se* with a theory that takes the values that enter into composition to be context-dependent. If the theory tells us how the values vary with context, *and the variation is linguistically controlled*, the theory will still be revelatory of the tacit knowledge employed in pairing expressions with their encoded meanings.

So a theory which allows for context-dependent inputs to composition can still count as genuinely semantic, at least so long as the variation is linguistically controlled. But if the variation is *not* linguistically controlled, then even if the theory tells us how the values that enter into composition vary with context, what it is revelatory of is not the tacit knowledge we're interested in, but the *interaction* of that knowledge and pragmatic processing. In short, such a theory won't be genuinely semantic, but will rather be the result of a blurring of the boundary between semantics and pragmatics. (I think this is one way of spelling out the worry that I hinted at right at the beginning of the lectures.)

Is this a decisive objection to Recanati's solution? Perhaps not. The theory itself will not directly reveal the tacit knowledge we're interested in *qua* semanticists, but Recanati might suggest that we can work backwards—given an account of how linguistically free processes work together with linguistically controlled processes to yield contextual values. If we want to take that approach, we're thrown back on to the issues we've been thinking about over the past few weeks: the nature of the word meanings and of the various different pragmatic processes that enter into interpretation. In effect, a compositional theory of truth conditions presents an equation with two unknowns, but maybe we can, as it were, solve for the semantic by controlling for the pragmatic.

b. The second worry relates to pragmatics. The worry is that, unless the values that enter into composition are linguistically controlled, the results of composition will be too unconstrained. In other words, our theory will wrongly predict readings to arise that *don't*—uses of (1), for example, to say that a given leaf is brown and valuable.

concept of heavy rain). That then enters into composition with the meanings of the other parts of the sentence to yield RAIN*$_{\text{present}}$. This is then modulated further, giving [RAIN*$_{\text{present}}$] IN LONDON.

4 Stanley's Challenge

This brings us to Stanley's challenge (see 2007). This a challenge for anyone who holds that language is not compositional—though I think it is also a challenge for those like Recanati who hold that is compositional, but only in the weak form described above.

The challenge begins with the observation that we *seem* to grasp what is said (explicatures) by decoding the linguistic meaning, working out the contextual values where required (indexicals, ambiguities), and combining the results in accordance with rules corresponding to syntactic composition. If Fodor and contextualists are correct, and language is not compositional, there is nothing to prevent what is said with a sentence on a particular occasion from containing constituents that do not correspond to any element in the sentence's logical form. Indeed, that's exactly what they claim. Consider:

(1) 'It's raining'
[$\text{RAIN}_{\text{present}}$] IN LONDON

But now, the contextualist seems to wrongly predict that what is said with the sentences in (a) could be as given in (b) — given only an appropriate context. For example, in (2), it might be contextually salient that everyone who likes Sally also likes their mother:

(3) (a) John likes Sally
(b) JOHN_1 LIKES SALLY AND HIS_1 MOTHER

(4) (a) Every Frenchman is seated
(b) EVERY FRENCHMAN OR DUTCHMAN IN THE ROOM IS SEATED.

Stanley concludes that language must be compositional. The composition rules apply to the contents of linguistic expressions, and there can be pragmatic contributions to what is said only where this is mandated by elements of logical form, e.g. by indexicals.

5 Indexicalism

To account for apparent violations of compositionality, Stanley posits extensive covert linguistic structure — **hidden indexicals**. At the level of logical form, the verb 'rain', for example, is of the form [RAIN (L)], where 'L' is a variable that takes a location, e.g. LONDON, as value in context. However, there is no part of (2a) to which the "constituent" AND HIS MOTHER could be assigned — so reading (2b) doesn't arise.

Similar treatments are given of other apparent counterexamples to compositionality. Szabó (2001) attempts to account for variation in what is said with 'This leaf is green', for example, by positing a variable ranging over *parts* in the adjective 'green'. Weiskopf (2007) attempts to account for compound nominals ('axe murderer'/'child murderer')

by positing a variable ranging over relations between the two parts of the compound.

But the "hidden indexical" strategy is problematic. For one thing, the different things that can be said with, e.g., 'green' seem to differ from one another in arbitrarily many ways. It's implausible to think all of the parameters of possible variation can be specified in the form of indexicals at the level of logical form. Indeed, we might think that it misconstrues the challenge that Context Shifting Arguments / Travis cases raise. Given how pervasive the phenomenon seems to be, why should we think that *any* attempt to specify the relevant parameters will succeed? In other words, why *shouldn't* we think that, given such a specification, it will still be possible to find two different contexts, alike with respect to all the specified parameters, which are such that the same sentence can be used to say something true in one context and something false in the other?

There is a lot of discussion of indexicalism. See, for example, Recanati (2004, Ch. 7), Martí (2006), Neale (2007), Collins (2007), and Hall (2008), and references therein.

Back Up Reading and Questions

Reading

Read Stanley (2000), which is reprinted as chapter 1 of his (2007).

Questions

1. Like colour terms, many other expressions seem to be able to be modulated in a range of ways. Consider what parameters would be required to explain the interpretations of the underlined expressions in the following examples.

   a) 'It's raining'

      uttered as an answer to each of the following questions:

      i. Is it still snowing?

      ii. Are you saying it's still drizzling?

      iii. Is it still pouring like this morning?

      iv. Can we go for a walk now?

      v. Why did you call a taxi?

   b) i. A: Where's John? B: He's working.

      ii. A: How can John afford these expensive holidays? B: He's working.

   How would an indexicalist account account for the range of interpretations?

References

Collins, J. (2007) 'Syntax, More or Less' *Mind* **116** (464) p. 805.

Hall, A. (2008) 'Free Enrichment or Hidden Indexicals?' *Mind & Language* **23** (4) pp. 426–456.

Martí, L. (2006) 'Unarticulated Constituents Revisited' *Linguistics and Philosophy* **29** (2) pp. 135–166.

Neale, S. (2007) 'On Location' in M. O'Rourke and C. Washington, eds., *Situating Semantics: Essays on the Philosophy of John Perry* (MIT Press) pp. 251–393.

Pagin, P. (2005) 'Compositionality and Context' in G. Preyer and G. Peter, eds., *Contextualism in Philosophy: Knowledge, Meaning, and Truth* (Oxford: Oxford University Press) pp. 303–348.

Recanati, F. (2004) *Literal Meaning* (Cambridge University Press).

——— (2012) 'Compositionality, Flexibility, and Context Dependence' in *The Oxford Handbook of Compositionality* (Oxford: Oxford University Press) pp. 175–199.

Stanley, J. (2000) 'Context and Logical Form' *Linguistics and Philosophy* **23** (4) pp. 391–434.

——— (2007) *Language in Context: Selected Essays* (Oxford University Press).

Szabó, Z. G. (2001) 'Adjectives in Context' in *Perspectives on Semantics, Pragmatics, and Discourse: a Festschrift for Ferenc Kiefer* (Amsterdam: John Benjamins Publishing Company) pp. 119–146.

Weiskopf, D. A. (2007) 'Compound Nominals, Context, and Compositionality' *Synthese* **156** (1) pp. 161–204.